

COMP 6721 Applied Artificial Intelligence (Fall 2023)

Lab Exercise #09: Knowledge Graphs, Part II

Question 1 Analyze and reason about RDF Schema class hierarchies and property inheritance in a theoretical context. This exercise focuses on your understanding of the concepts without the need to write RDF triples or code.

- a) Describe the relationships and reasoning within a class hierarchy that includes **Person**, **Employee**, **Manager**, and **Intern**, where **Manager** and **Intern** are subclasses of **Employee**, and **Employee** is a subclass of **Person**.
- b) Discuss the implications of assigning a property **worksAt**, with a domain of **Employee** and a range of **Organization**, to an individual of type **Intern** in terms of property inheritance in RDFS.
- c) Conceptualize an advanced RDF Schema with a complex class hierarchy involving **Person**, **Author**, **Academic**, **Student**, **GraduateStudent**, and so on, and discuss the inferencing capabilities offered by such a schema.

Question 2 We'll now continue the knowledge graph exercise from the previous week. Your next task is to develop a new *vocabulary* using RDF Schema (RDFS): The *Friends-of-Concordia-University (FOCU)* schema. This vocabulary should be able to express:

- (i) Classes for **Student** and **Professor**, both of which are subclasses of **foaf:Person**
- (ii) A **University** class, which is a subclass of **foaf:Organization**
- (iii) A property describing that a **Student** is enrolled at a **University**
- (iv) A property describing that a **Professor** teaches **Students**

You can start with the following template (in Turtle format):

```
@prefix rdfs: <http://www.w3.org/2000/01/rdf-schema#> .
@prefix rdf: <http://www.w3.org/1999/02/22-rdf-syntax-ns#> .
@prefix xsd: <http://www.w3.org/2001/XMLSchema#> .
@prefix foaf: <http://xmlns.com/foaf/0.1/> .
@prefix focu: <http://focu.io/schema#> .
@prefix focudata: <http://focu.io/data#> .
```

```
# RDF Schema for the FOCU Vocabulary
```

```
focu:[CLASSNAME]
  a rdfs:Class ;
```

```
rdfs:label "[LABEL]"@en .
```

```
focu:[OTHERCLASSNAME]  
  a rdfs:Class ;  
  rdfs:subClassOf foaf:[CLASS] ;  
  rdfs:label "[LABEL]"@en .
```

```
focu:[PROPERTY]  
  a rdf:Property ;  
  rdfs:label "[LABEL]"@en ;  
  rdfs:comment "[COMMENT]"@en ;  
  rdfs:domain uni:[DOMAIN_CLASS] ;  
  rdfs:range uni:[RANGE_CLASS] .
```

- (a) Add your RDFS triples for the four definitions above by filling in the missing parts in [square brackets].
- (b) Validate your new RDF schema using the tools mentioned in last week's lab and visualize them in form of a graph.
- (c) Now add some triples using your new FOCU vocabulary that describe yourself as a student at Concordia. Add information about yourself (age, email) using the FOAF vocabulary as we did on the lecture Worksheet #8.
- (d) Again, validate and visualize your triples.

Some notes:

- There are of course existing vocabularies for universities etc. and in a real system you would not create a new vocabulary duplicating existing definitions (just like we re-use FOAF to describe a person here). This is just for exercise purposes, so that you see how to define a minimal, working RDF Schema.
- We are using `focu.io` as the authority (domain name): This is again bad practice, since you do not control this domain name and thus cannot publish anything using this authority.
- The template above defines `focudata:` as the namespace for the triples describing university people; you could use another namespace here, like the `example.org` we've used before. The point here is that it is not good practice to mix schema information (in this case your FOCU schema) with data described using this schema into the same namespace. In other words, you want to be able to load the schema and the data separately from each other into a system.
- Can you add a triple that relates your new `University` definition to an existing concept in the FOAF vocabulary?

Question 3 Using your RDFLib program from the previous week, read the FOAF+FOCU triples about some people you created above. Next, your program should ask for a name and retrieve all the available information from the graph. For example, given the name Joe, it will print out all available triples like age, email, etc.:

```
import rdflib
from rdflib import Literal, Namespace
from rdflib.namespace import FOAF
```

Add the user-defined namespaces:

```
# Create the user-defined namespaces
FOCU = Namespace("http://focu.io/schema#")
FOCUDATA = Namespace("http://focu.io/data#")
```

Here is the command to create a graph instance:

```
g = rdflib.Graph()
```

Now parse the graph containing your triples created in Question 1:

```
g.parse("Lab09.ttl", format="turtle")
```

Use the command below to start accepting user input:

```
user_input = input('Hello, I am your smart university agent. Who are you
looking for? ')
```

```
Hello, I am your smart university agent. Who are you looking for?
> Joe
```

```
Joe is 22 years old and has the email address joe@example.com.
```

Implement searching and printing the information by going through the graph.

Hint: Use the RDFLib functions you learned in the previous lab to navigate the graph and retrieve the required information, i.e., age and email address. Here, the user inputs a person's first name, therefore find the triples matching the given first name and then proceed to retrieve the required information from the graph.

Question 4 Next, if the person is enrolled at a university, your agent should determine the *location* of the university (using the triples about the university from DBpedia) and print out this information as well:

```
Hello, I am your smart university agent. Who are you looking for?  
> Jane
```

```
Jane is 24 years old and studies at Concordia University,  
which is located in Montreal, Quebec, Canada
```

This time, we'll use DBpedia to retrieve additional triples, so make sure the student in your knowledge base is linked to, e.g., http://dbpedia.org/resource/Concordia_University.

Hint: The only new information we need here is the location. To obtain it, you can merge the DBpedia graph about the university into your knowledge base and fetch all the additional information from there.

Use the following namespace declaration:

```
DBP = Namespace("http://dbpedia.org/property/")
```

and find the correct property to use for the location in DBpedia.

You can print out additional information about the university if you like.

Question 5 We'll now make the AI even smarter, by also supporting the query "What is <X>", where you retrieve triples about <X> from DBpedia, printing out the abstract information, e.g.:

```
Hello, I am your smart university agent. How can I help you?  
> What is Concordia University?
```

```
Concordia University (French: Université Concordia) is a public  
research university located in Montreal, Quebec, Canada. Founded in  
1974 following the merger of Loyola College and Sir George Williams  
University, Concordia is one of the three universities in Quebec  
where English is the primary language of instruction (the others  
being McGill and Bishop's). . . .
```

To implement this, we'll start by checking whether a "What is <X>" pattern is present in the user's input:

```
# Define pattern  
pattern = "What is"  
  
if pattern in user_input:
```

```
# Extract the entity name after the pattern,  
# handling potential trailing question mark  
entity_name = user_input.split(pattern, 1)[1].split("?", 1)[0].strip()
```

Implement the following two steps:

- (a) Create a DBpedia URI based on the input subject `<X>` and load the graph into your program. For example, if `<X>` is `Concordia University`, create the URI http://dbpedia.org/resource/Concordia_University.¹
- (b) Use the RDFLib functions as before to navigate and retrieve the required output from the graph.² You'll need to find the right property to use; generally, you can find an *abstract* for most entries in DBpedia.

Notice how similar this exercise is to the functionality of AI assistants like Google Assistant, Siri, or Alexa: By querying DBpedia for 'What is `<X>?`', you're emulating their technique of using knowledge graphs to provide informative answers. This process, known as *grounding*, is a departure from the pattern-matching approach seen in earlier chatbots like Eliza, demonstrating a more advanced and context-aware method of responding to user queries.

In addition to powering AI assistants, knowledge graphs have significant applications across various industries. For instance, they are instrumental in health-care for organizing complex medical data, aiding in diagnosis and patient care. In the financial sector, knowledge graphs facilitate fraud detection and risk management by analyzing intricate transaction networks. While distinct from Large Language Models (LLMs) like GPT-3, which generate responses based on pattern recognition in extensive text data, knowledge graphs offer a structured, queryable approach to handling information. This structured approach is particularly valuable in scenarios requiring precise, verifiable data retrieval or where relationships between different entities are crucial. Understanding knowledge graphs thus provides foundational insights into a key aspect of AI, relevant for fields ranging from semantic search and recommendation systems to complex data analysis in research and industry.

¹Of course, this naïve approach of converting a string to a DBpedia URI does not always work in practice, but there are more sophisticated ways available of achieving this, for example, the DBpedia Spotlight tool: <https://www.dbpedia-spotlight.org/>.

²Another way of implementing this would be to generate SPARQL queries, similar to how you've done it on Worksheet #8. However, we're not covering SPARQL in this lab.